



# GPU Architecture

National Tsing-Hua University  
2019, Summer Semester

# Outline

## ■ Thread execution

- Execution model
- Warp
- Warp Divergence


## ■ Memory hierarchy

# Execution Model

## Software



Thread



Thread block



Grid

## Hardware



Scalar processor



Stream Processor (SM)



GPU device

Threads are executed by scalar processor

Thread blocks are executed on SM  
Several concurrent thread block can reside on one SM

A kernel is launched as a grid of thread blocks

# Thread Execution

- CUDA threads are grouped into blocks
  - All threads of the same block are **executed in an SM**
  - SMs have **shared memories**, where threads within a block **can communicate**
  - **The entire threads of a block must be executed completely before there is space to schedule another thread block**
- Hardware schedules thread blocks onto available SMs
  - **No guarantee of order of execution**
  - If an SM has more resources, the hardware can schedule more blocks

# Warp

- Inside the SM, threads are launched in groups of 32, called **warps**

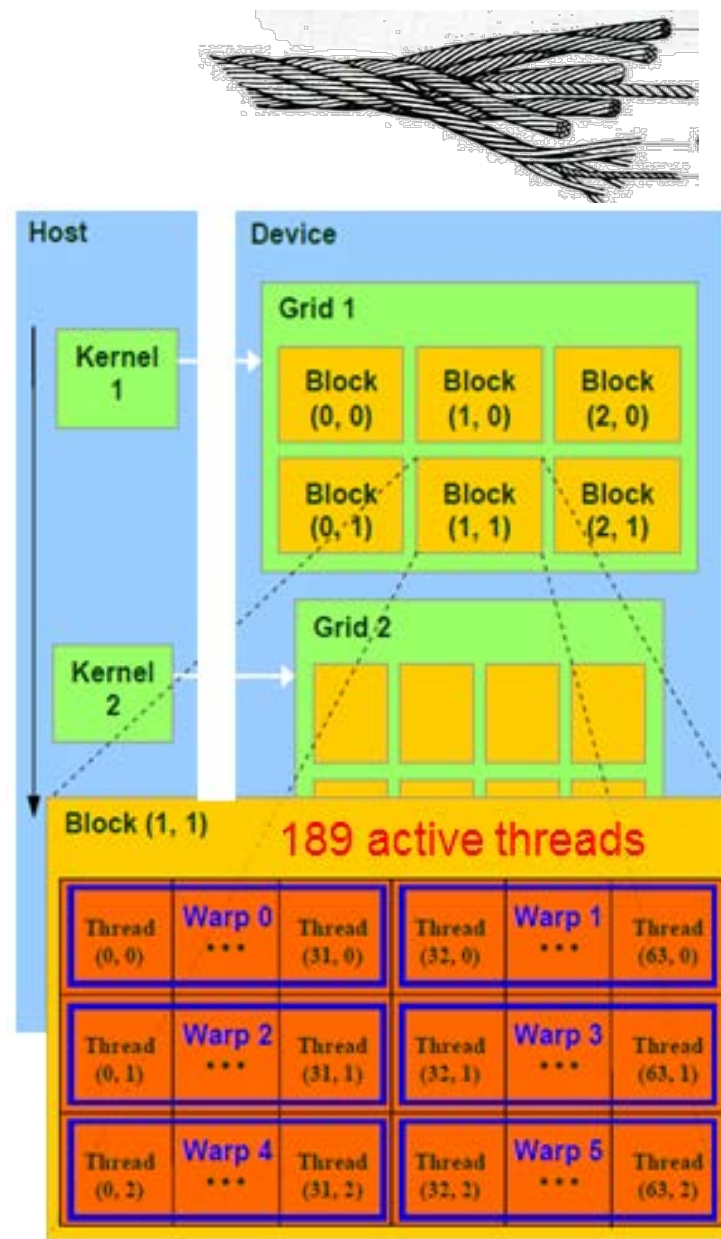
- Warps share the control part (**warp scheduler**)
- At any time, **only one warp** is executed per SM
- Threads in a warp will be executing the same instruction (SIMD)

- In other words ...

- Threads in a warp execute **physically** in parallel
- Warps and blocks execute **logically** in parallel
- ➔ Kernel needs to sync threads within a block

- For Fermi:

- Maximum number of active blocks per SM is 8
- Maximum number of active warps per SM is 48
- Maximum number of active threads per SM is  $48 \times 32 = 1,536$

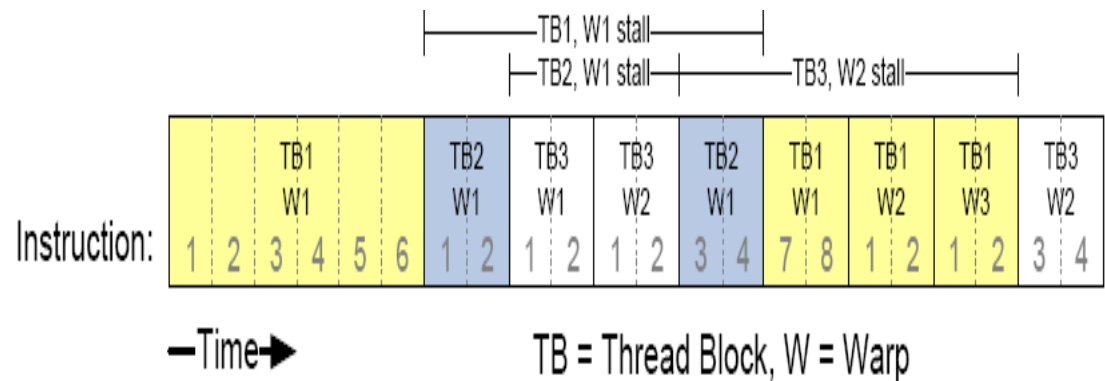
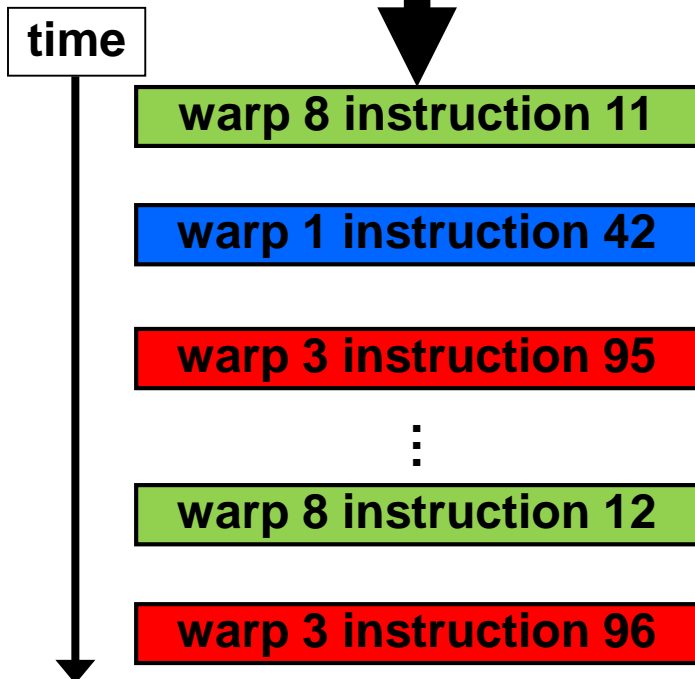


# Warp Scheduler



SM multithreaded  
Warp scheduler

- SM hardware implements zero-overhead Warp scheduling
  - Warps whose next instruction has its **operands ready for consumption** are **eligible for execution**
  - Warps are **switched when memory stalls**
  - Eligible Warps are selected for execution on **prioritized scheduling**
  - **All threads in a Warp execute the same instruction when selected**



# Warp Divergence

- What if different threads **in a warp** need to do different things:


- Including any flow control instruction (**if, switch, do, for, while**)

```
if (foo(threadIdx.x)) {  
    do_A();  
} else {  
    do_B();  
}
```

- Different execution paths within a warp are serialized

- Predicated instructions which are carried out only if logical flag is true
- All threads compute the logical predicate and two predicated instructions/statements

➔ **Potential large lost of performance**



# Avoid Diverging in a Warp

## ■ Example with divergence:

```
if (threadIdx.x > 2) { ... }  
else { ... }
```

- Branch granularity < warp size

## ■ Example without divergence:

```
if (threadIdx.x / WARP_SIZE > 2) { ... }  
else { ... }
```

- Different warps can execute different code with no impact on performance
- Branch granularity is a whole multiple of warp size



# Iteration Divergence

- A single thread can drag a whole warp with it for a long time
- Know your data patterns
- If data is unpredictable, try to flatten peaks by letting threads work on multiple data items

# Unroll the for-loop

- Unroll the statements can reduce the branches and increase the pipeline
- Example:

```
for (i=0;i<n;i++) {  
    a = a + i;  
}
```

➤ Unrolled 3 times

```
for (i=0;i<n;i+=3) {  
    a = a + i;  
    a = a + i+1;  
    a = a + i+2;  
}
```

# #pragma unroll

- The `#pragma unroll` directive can be used to control unrolling of any given loop.
- must be placed immediately before the loop and only applies to that loop
- Example:  

```
#pragma unroll 5  
for (int i = 0; i < n; ++i)
```

  - the loop will be **unrolled 5 times**.
  - The compiler will also insert code to ensure correctness
- `#pragma unroll 1` will prevent the compiler from ever unrolling a loop.

# Atomic Operations

- Occasionally, an application may need threads to update a counter in **shared** or **global** memory

```
__shared__ int count;
```

```
.....
```

```
if (.....) count++;
```

- Synchronization problem: if two (or more) threads execute this statement at the same time
- Solution: use atomic instructions supported by GPU
  - addition / subtraction
  - max / min
  - increment / decrement
  - compare-and-swap

# Example: Histogram

```
/* Determine frequency of colors in a picture  
colors have already been converted into ints. Each  
thread looks at one pixel and increments a counter  
atomically*/
```

```
__global__ void hist(int* color, int* bin){  
    int i = threadIdx.x + blockDim.x *  
            blockDim.x;  
  
    int c = colors[i];  
    atomicAdd(&bin[c], 1);  
}
```

# Example: Global Min/Max

```
__global__ void global_max(int* values,  
int* gl_max){  
    int i = threadIdx.x + blockDim.x *  
                                blockDim.x;  
    int val = values[i];  
    atomicMax(gl_max, val);  
}
```

- Not very fast for data in shared memory
- Only slightly slower for data in device memory

# Outline

- Thread execution
- Memory hierarchy
  - Register & Local memory
  - Shared memory
  - Global & Constant memory

# GPU Memory Hierarchy

## ■ Registers

- Read/write per-thread
- Low latency & High BW

## ■ Shared memory

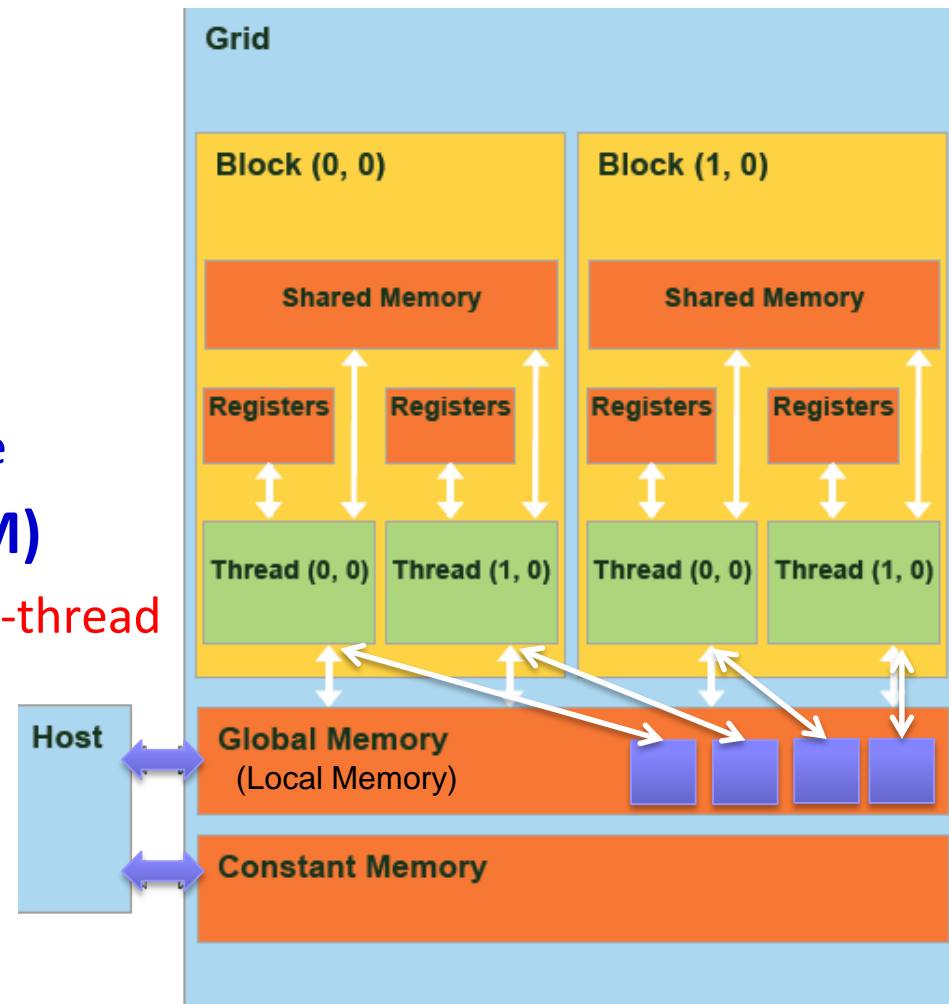
- Read/write per-block
- Similar to register performance

## ■ Global/Local memory (DRAM)

- Global is per-grid & Local is per-thread
- High latency & Low BW
- Not cached

## ■ Constant memory

- Read only per-grid
- Cached





# Memory Access

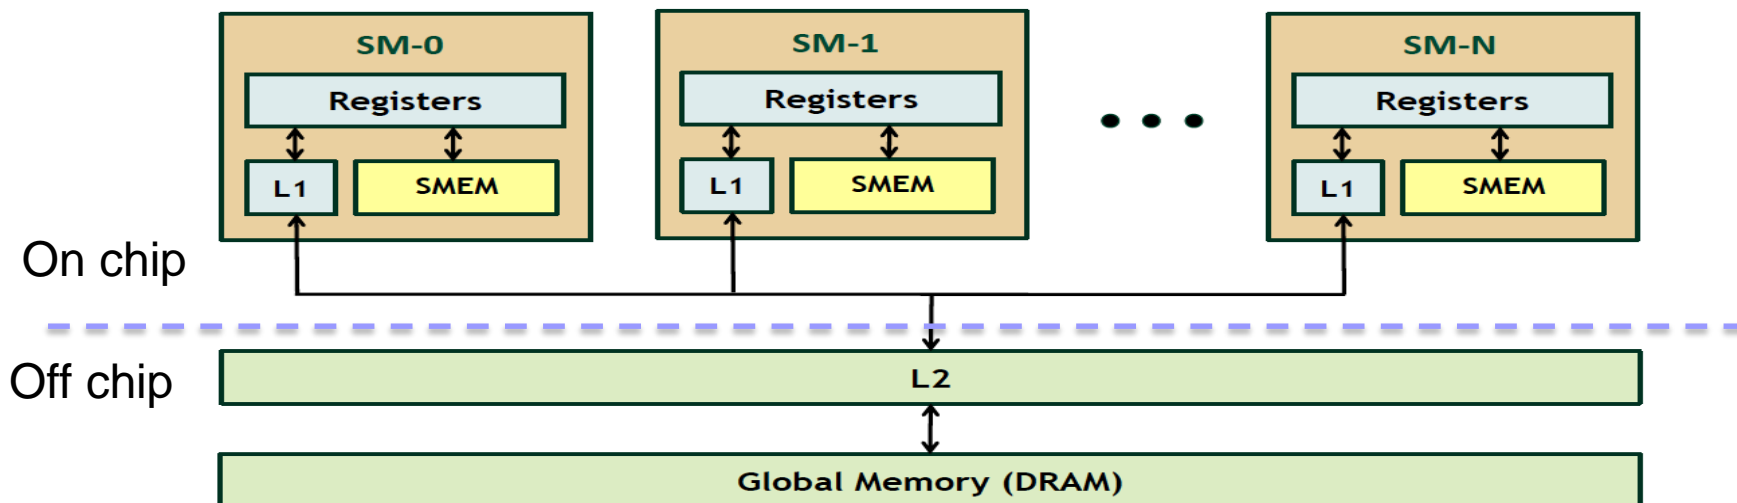
## ■ A store writes a line to L1

- If evicted, that line is written to L2
- The line could also be evicted from L2 and written to DRAM (global mem.)

## ■ A load requests the line from L1

- If a hit, operation is complete
- If a miss, then requests the line from L2
  - ◆ If a miss, then requests the line from DRAM (global memory)

## ■ Only GPU threads can access local memory addresses

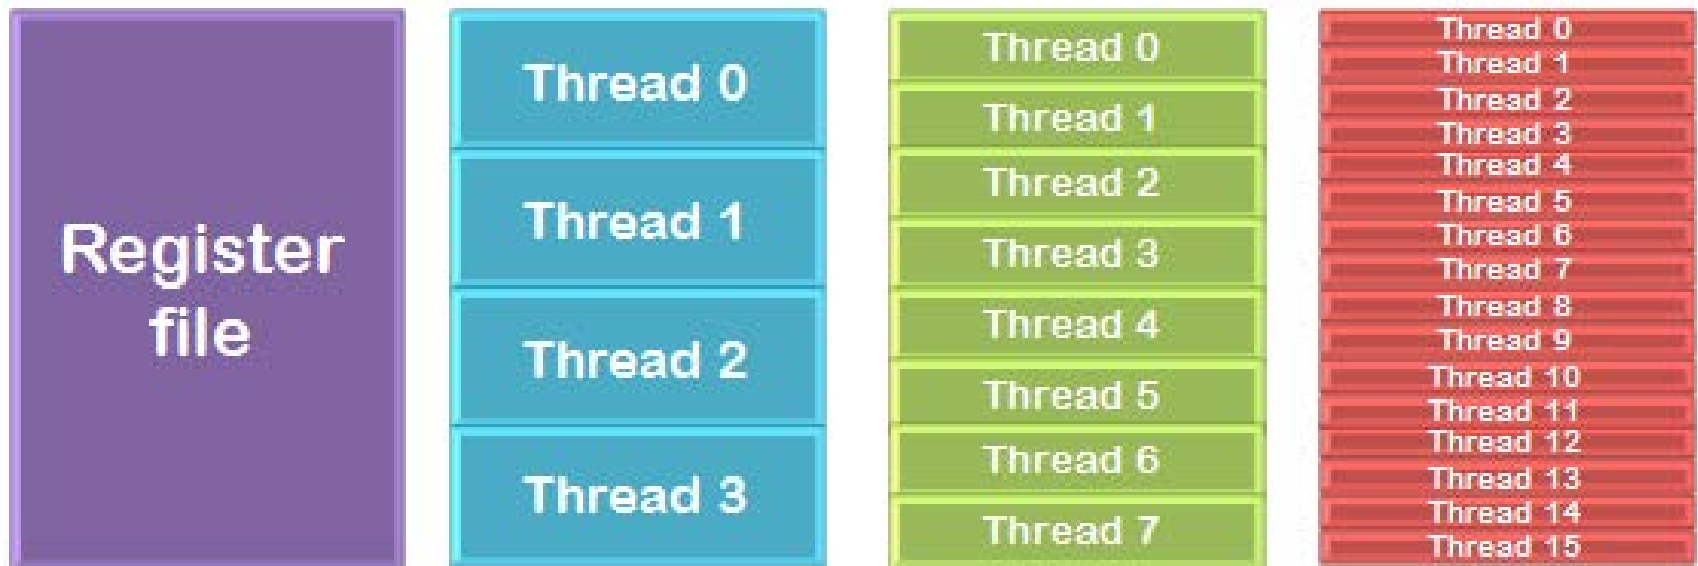


# Register

- Register consumes zero extra clock cycles per instruction, except
  - Register read-after-write dependencies (24 cycles) and
  - Register memory bank conflicts
- Register spilling
  - Max number of register per threads is 63
  - Local memory is used if the register limit is met
  - **Array variables** always are allocated in **local memory (DRAM)**
  - Max amount of local memory per thread is 512K

# Register Pressure

- Too few threads
  - can't hide pipeline / memory access latency
- Too many threads
  - register pressure
  - Limited number of registers among concurrent threads
  - Limited shared memory among concurrent blocks



# Local Memory

- Name refers to memory where registers and other thread-data is spilled
  - Usually when one runs out of SM resources
  - “**Local**” because each thread has its own private area
- Details:
  - Not really a “memory” – bytes are stored in **global memory (DRAM)**
- Differences from global memory:
  - Addressing is resolved by the **compiler**
  - Stores are cached in L1

# Example

```
__device__ void distance(int m, int n, int *V){  
    int i, j, k;  
    int a[10], b[10], c[10];  
    ...  
}
```

- Variables *i, j, k, a, b, c* are called “local variables”.
- It is likely that variable *i, j, k* are stored in registers, and variable *a, b, c* are stored in “local memory” (off-chip DRAM).
  - Compiler decides which memory space to use.
  - Registers aren’t indexable, so arrays have to be placed in local memory.
  - If not enough registers, local memory will be used.
- Only allowed static array!! ➔ No `int a[m];`

# Outline

- Thread execution
- **Memory hierarchy**
  - Register & Local memory
  - **Shared memory**
  - Global & Constant memory
- Occupancy

# Shared Memory

- Programmable cache!!

- **Almost as fast as registers**

- Scope: shared by all the threads in a block.

- The threads in **the same block** can communicate with each other through the **shared memory**.

- Threads in **different blocks** can only communicate with each other through **global memory**.


- Size: at most 48K per block

- On Fermi/Kepler GPU, **shared memory and L1 cache use the same memory hardware (64K)**. Programmers can decide the ratio of shared memory and L1 cache:

- The ratio (shared:L1) can be (3:1) or (1:1) or (1:3).

# General Strategy

1. Load data from global memory to shared memory
2. Process data in the shared memory
3. Write data back from shared memory to





# APSP Parallel Implementation Revisit

- Use  **$n*n$**  threads.
- Each updates the shortest path of one pair vertices
- Use **global memory** to store the matrix D.

```
__global__ void FW_APSP(int k, int D[n][n]) {  
    int i = threadIdx.x;  
    int j = threadIdx.y;  
    if (D[i][j] > D[i][k] + D[k][j])  
        D[i][j] = D[i][k] + D[k][j];  
}  
  
int main() { ...  
    dim3 threadsPerBlock(n, n);  
    for (int k = 0; k < n; k++)  
        FW_APSP<<<1, threadsPerBlock >>>(k, D);  
}
```

6 global  
memory  
access

# Using Shared Memory

- This way of using shared memory is called **dynamic allocation of shared memory**, whose size is specified in the kernel launcher.

```
FW_APSP<<<1, n*n, n*n*sizeof(int)>>> (...);
```

- The third parameter is the size of shared memory.

```
extern __shared__ int S[][];  
__global__ void FW_APSP(int k, int D[n][n]) {  
    int i = threadIdx.x;  
    int j = threadIdx.y;  
    S[i][j]=D[i][j]; // move data to shared memory  
    __syncthreads();  
    // do computation  
    if (S[i][j]>S[i][k]+S[k][j])  
        D[i][j]= S[i][k]+S[k][j];  
}
```

**ONLY 2**  
global mem  
access

# Limit of Dynamic Allocation

- If you have multiple extern declaration of shared:

```
extern __shared__ float As[];
```

```
extern __shared__ float Bs[];
```

this will lead to As pointing to the same address as Bs.

- Solution: keep As and Bs inside the 1D-array.

```
extern __shared__ float smem[];
```

- Need to do the memory management yourself

- When calling kernel, launch it with size of  $sAs + sBs$ , where  $sAs$  and  $sBs$  are the size of As and Bs respectively.
- When indexing elements in As, use `smem[ 0 :  $sAs - 1$  ]` ;  
when indexing elements in Bs, use `smem[  $sAs$  :  $sAs + sBs$  ]`.

# Using Shared Memory

- `FW_APSP<<<1, n*n, n*n*sizeof(int)>>> (...);`
  - The third parameter is the size of shared memory.

```
extern __shared__ int S[];
__global__ void FW_APSP(int* k, int* D, int* n) {
    int i = threadIdx.x;
    int j = threadIdx.y;
    S[i*(*n)+j]=D[i*(*n)+j]; //move data to shared memory
    __syncthreads();
    // do computation
    if (S[i*(*n)+j]>S[i*(*n)+k]+S[k*(*n)+j])
        D[i*(*n)+j]= S[i*(*n)+k]+S[k*(*n)+j];
}
```

# Static Shared Memory Allocation

- If the size of shared memory is known in compilation time, shared memory can be allocated statically.

```
__global__ void FW_APSP(int k, int D[][]){  
    __shared__ int DS[10*10];  
    ...  
}
```

Must know  
n=10 at  
compile time

# Outline

- Thread execution
- **Memory hierarchy**
  - Register & Local memory
  - Shared memory
  - **Global & Constant memory**

# Global Memory in Kernel

- Through the kernel launcher arguments
  - Need to use `cudaMalloc` to allocate memory and use `cudaMemcpy` to set their values.
  - This method is what we used in previous examples.


```
cudaMemcpy( void *dst, const void  
*src, size_t count, enum  
cudaMemcpyKind kind)
```

# How to Allocate Device (Global) Memory

1. `cudaMalloc(void **devPtr, size_t size)`

- `devPtr`: return the address of the allocated memory on device
- `size`: the allocated memory size (bytes)

2. `cudaFree (void *devPtr)`





# Constant Memory

- Same usage and scope as the global memory except
  - **Read only**
  - Declare by variable qualifier `__constant__`
  - Move by `cudaMemcpyToSymbol()` & `cudaMemcpyFromSymbol()`
- Each SM has its own constant memory
  - The constant memory on each SM is of size 64K, and has a separated cache, of size 8K.

```
__constant__ int constData[100];
int main(void) {
    int A[100];
    cudaMemcpyToSymbol(constData, A, sizeof(A));
    add<<<grid_size,blk_size>>>();
    cudaMemcpyFromSymbol(A, constData, sizeof(A));
}
__global__ kernel() {
    int v = constData[threadIdx];
}
```

# CUDA Variables within a Kernel

Variable declaration	Memory	Scope	Lifetime
<code>int var</code>	Register	Thread	Thread
<code>int array_var[10]</code>	Local	Thread	Thread
<code>__shared__ int shared_var</code>	Shared	Block	Block
<code>__device__ int global_var</code>	Global	Grid	App
<code>__constant__ int constant_var</code>	Constant	Grid	App

- Scalar variables without qualifier reside in a register
  - Compiler will spill to thread local memory
- Array variables without qualifier reside in thread-local memory

# Memory Speed

Variable declaration	Memory	Speed
<code>int var</code>	Register	1x
<code>int array_var[10]</code>	Local	100x
<code>__shared__ int shared_var</code>	Shared	1x
<code>__device__ int global_var</code>	Global	100x
<code>__constant__ int constant_var</code>	Constant	1x

- Scalar variables reside in fast, on-chip registers
- Shared variables reside in fast, on-chip memories
- Thread-local arrays & global variables reside in uncached off-chip memory
- Constant variables reside in cached off-chip memory

# Memory Scale

Variable declaration	Total no. of variables	Visible by no. of threads
<code>int var</code>	100,000	1
<code>int array_var[10]</code>	100,000	1
<code>__shared__ int shared_var</code>	100	100
<code>__device__ int global_var</code>	1	100,000
<code>__constant__ int constant_var</code>	1	100,000

- 100Ks per-thread variables, R/W by 1 thread
- 100s shared variables, each R/W by 100s of threads
- Global variable is R/W by 100Ks threads
- 1 constant variable is readable by 100Ks threads

# Quiz

## CUDA Variables within a Kernel

Local  
variable

Shared  
variable

Global  
variable

Register

Local  
Memory

Shared  
Memory

Global  
Memory


Constant  
Memory

## GPU Memory Hierarchy

Q: What is the data scope of each variable, and what is the mapping between software and hardware?

# Quiz

## CUDA Variables within a Kernel



## GPU Memory Hierarchy

Q: What is the data scope of each variable, and what is the mapping between software and hardware?

# Reference

## ■ **NVIDIA CUDA Library Documentation**

- [http://developer.download.nvidia.com/compute/cuda/4\\_1/rel/toolkit/docs/online/index.html](http://developer.download.nvidia.com/compute/cuda/4_1/rel/toolkit/docs/online/index.html)

## ■ **NVIDIA CUDA Warps and Occupancy**

- [http://on-demand.gputechconf.com/gtc-express/2011/presentations/cuda\\_webinars\\_WarpsAndOccupancy.pdf](http://on-demand.gputechconf.com/gtc-express/2011/presentations/cuda_webinars_WarpsAndOccupancy.pdf)

## ■ **Heterogeneous computing course slides from Prof. Che-Rung Lee**